
The background of the slide is a photograph of a high-altitude mountain landscape. A large, jagged mountain peak is visible on the left, partially covered in snow. In the foreground, a wide, snow-covered slope leads up towards the mountain. Several climbers are visible on this slope, moving upwards. They are wearing outdoor gear, including backpacks and using ice axes. The sky is a clear, deep blue. The overall scene conveys a sense of challenge and achievement in a harsh environment.

Seminar Software Analytics Review of Automatically Learning Semantic Features for Defect Prediction

Presentation Meeting
20.01.2017

Felix Schober

The background of the slide is a photograph of a high-altitude mountain landscape. In the foreground, a snow-covered ridge leads towards a series of jagged, snow-dusted mountain peaks under a clear blue sky. Several hikers are visible on the ridge, moving away from the viewer. The overall scene is bright and crisp, with high contrast between the white snow and the blue sky.

Seminar Software Analytics Review of

Why Dreaming Networks Can Help a Company to Save Costs


Presentation Meeting
20.01.2017

Felix Schober

Agenda




Defect Prediction - Introduction



Technical Background



Discussion of Results



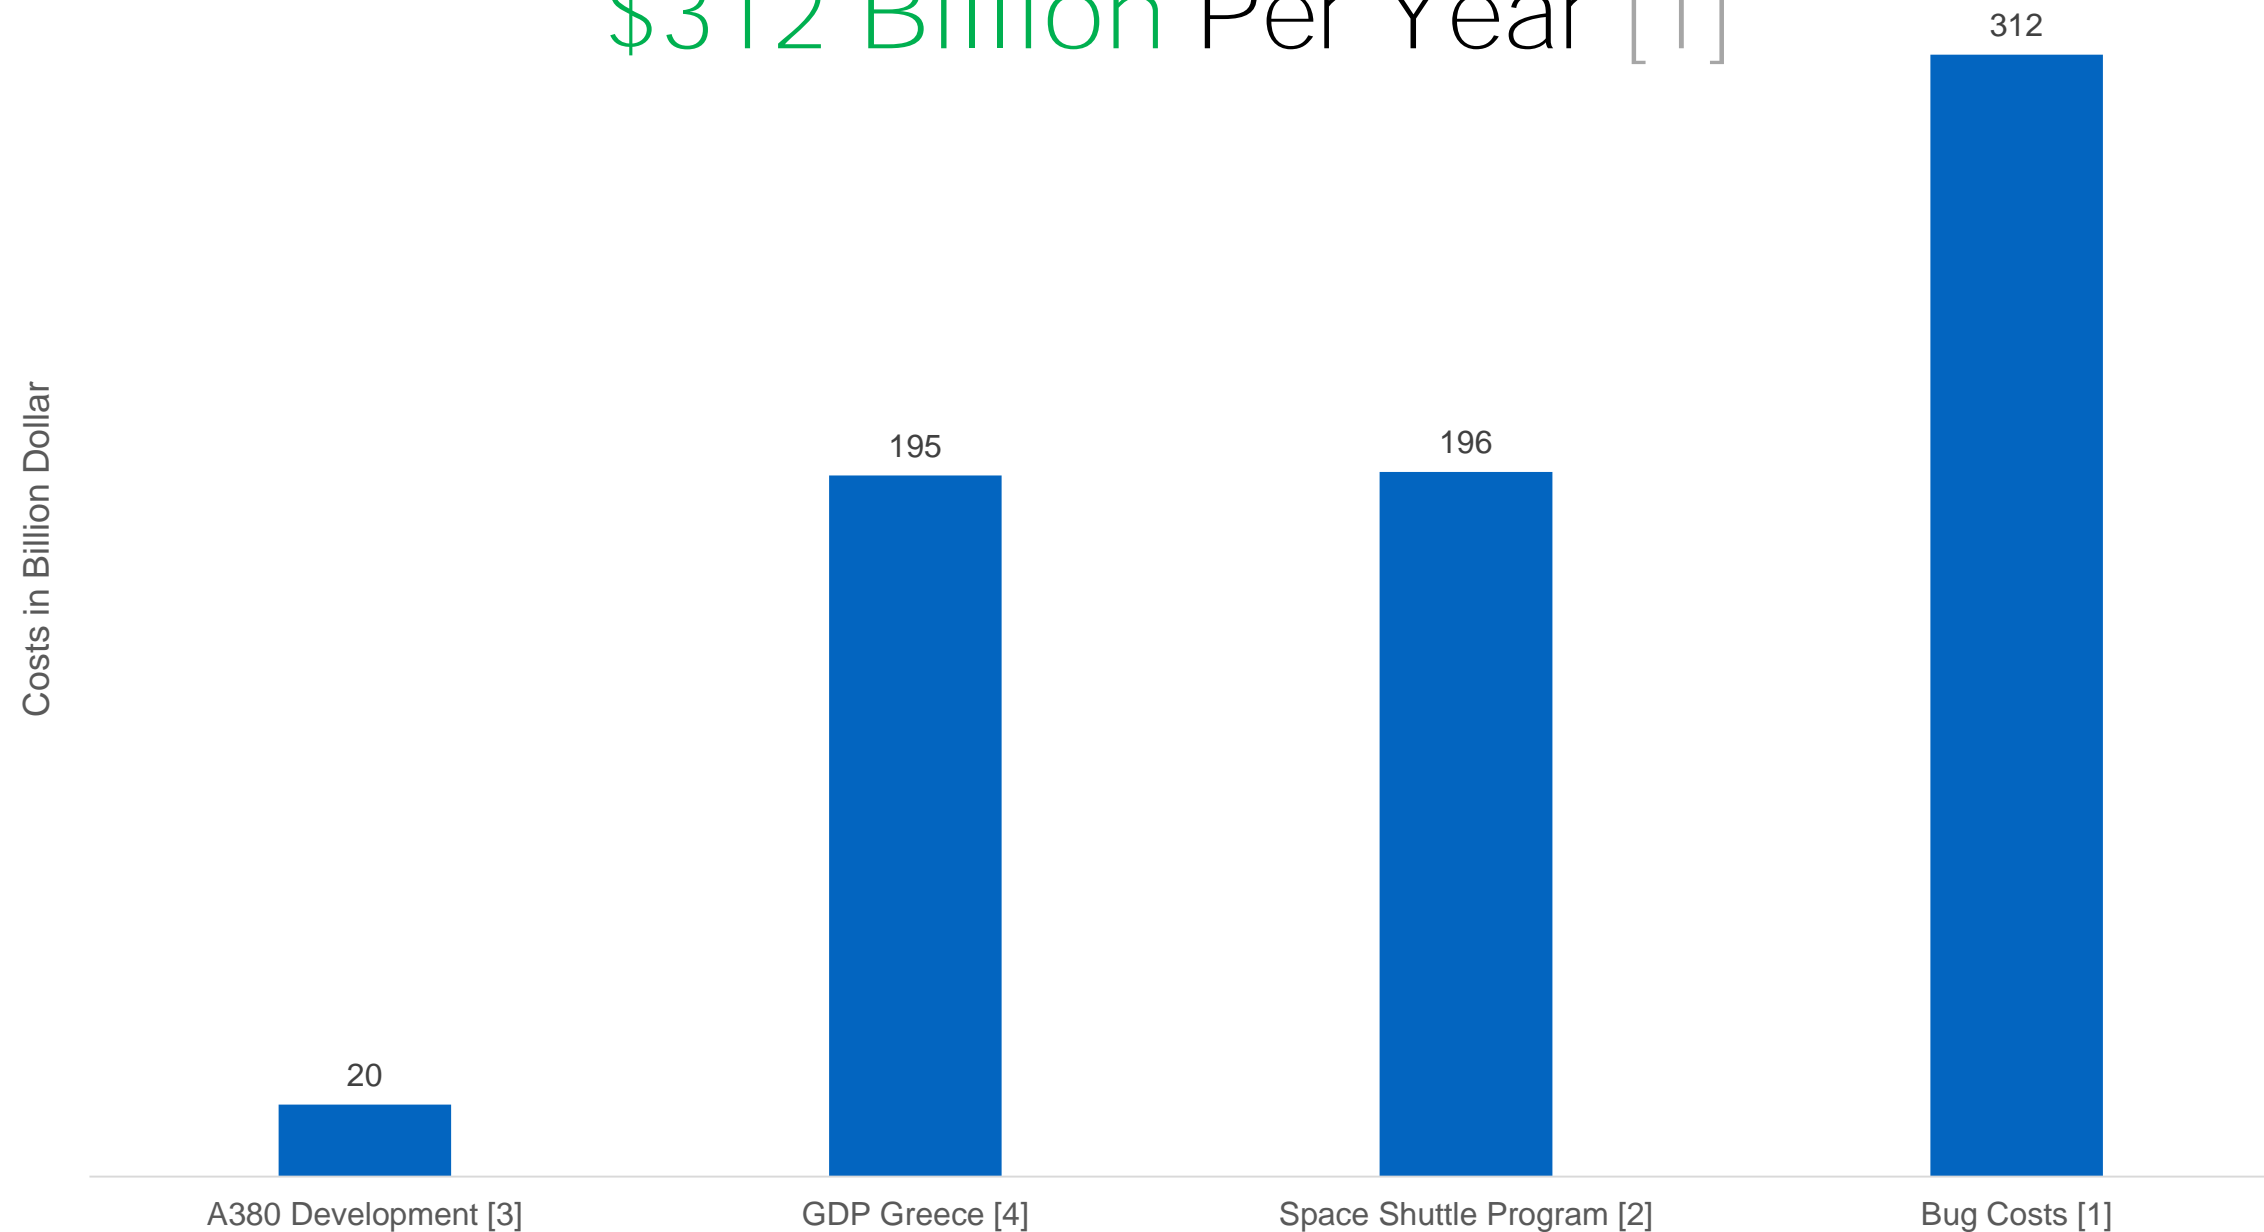
Demo



Conclusion

Cost of Defects

Global economy cost for bugs / software defects:
\$312 Billion Per Year [1]



[1] F. Brady "Cambridge University Study States Software Bugs Cost Economy \$312 Billion Per Year" prweb (2013) accessed 12.01.2017
<http://www.prweb.com/releases/2013/1/prweb10298185.htm>

[2] R. Pielke, R. Byerly "Shuttle Programme Lifetime Cost." In Nature Vol. 472 (2011)

[3] "Milliardengrab A380" Handelsblatt (2012) <http://www.genios.de/presse-archiv/artikel/HB/20121108/milliardengrab-a380/2B70894B-6C35-4D7B-A956-B0B5783C418D.html>

[4] World Bank – Last accessed 12.01.2017

http://data.worldbank.org/indicator/NY.GDP.MKTP.CD?locations=GR&name_desc=true

Defects in Software

“The application division at Microsoft experiences about 10 - 20 defects per 1000 lines of code during in-house testing, and 0.5 defect per 1000 lines of code in released product [...]” [5]

Windows Vista contains 50 Million Lines of code [6]

~25,000 Defects
in the released product

[5] S. McConnell “Code Complete 2nd edition.” Microsoft Press (2004) page 521
[6] S. Lohr, J. Markoff “Windows Is So Slow, but Why?” New York Times (2006) Last accessed 12.01.2017
<http://www.nytimes.com/2006/03/27/technology/27soft.html?adxnnl=1&pagewanted=all&adxnnlx=1382805118-0jnNRGXEVp3xoW+BDp8Q>

Defect Prediction

25,000 Defects
in the released product


50x – 200x
more expensive to fix a bug at a later stage [7]

Automatically detect defects as quick as possible

Agenda




Defect Prediction



Technical Background



Discussion of Results

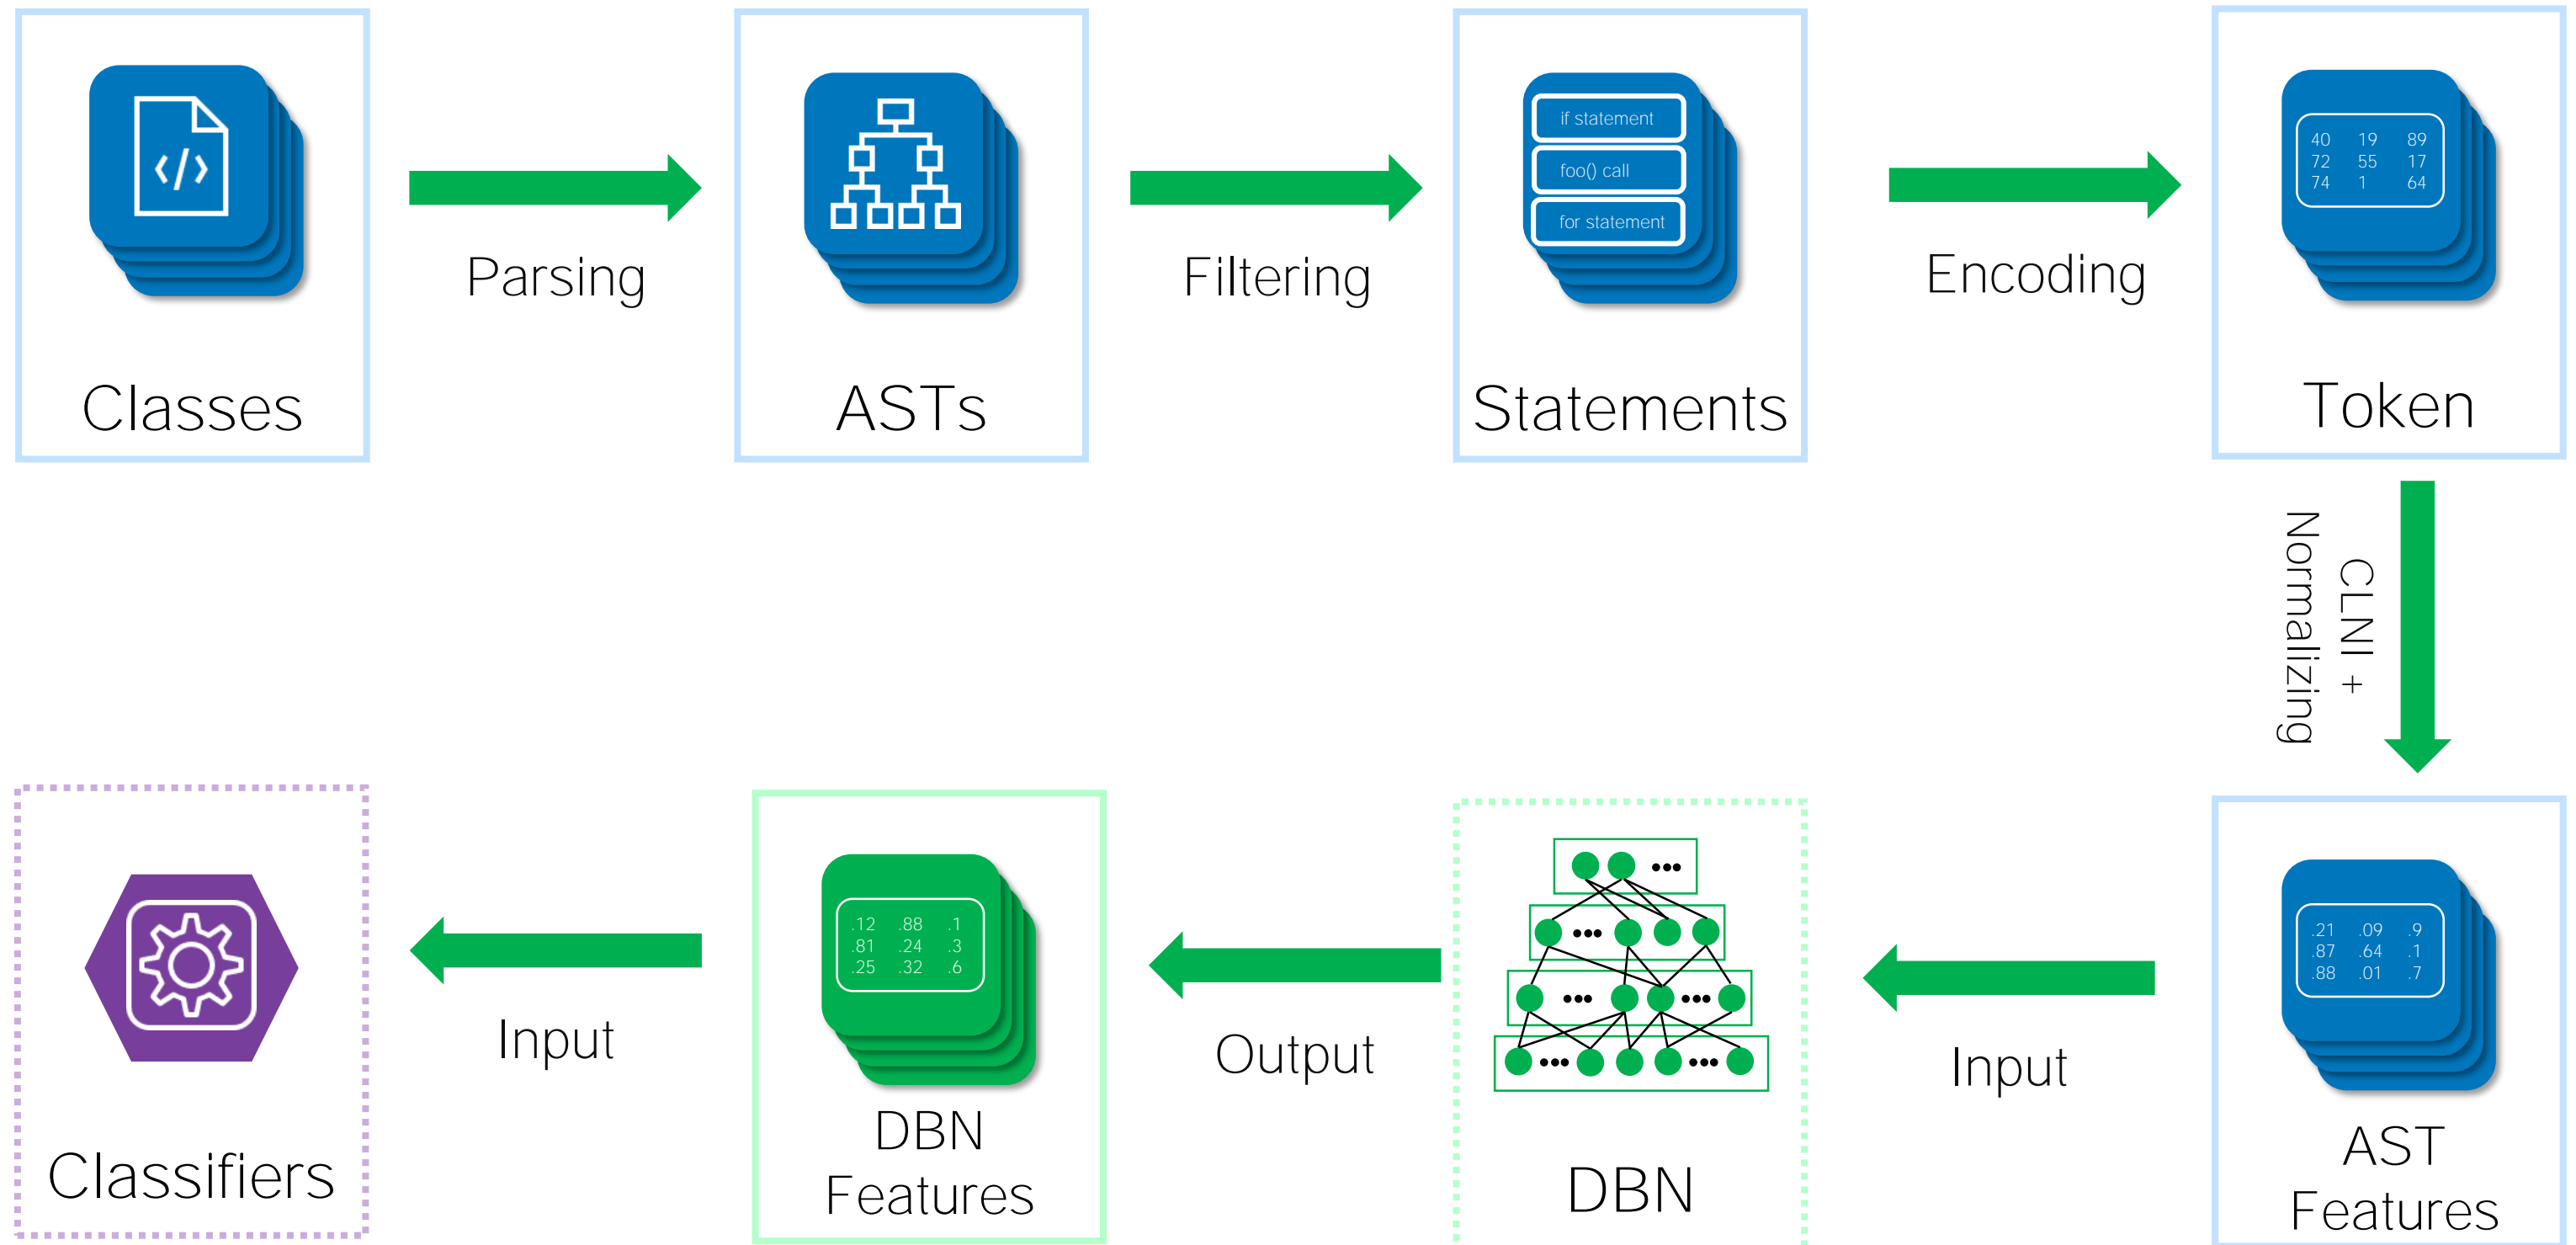


Demo

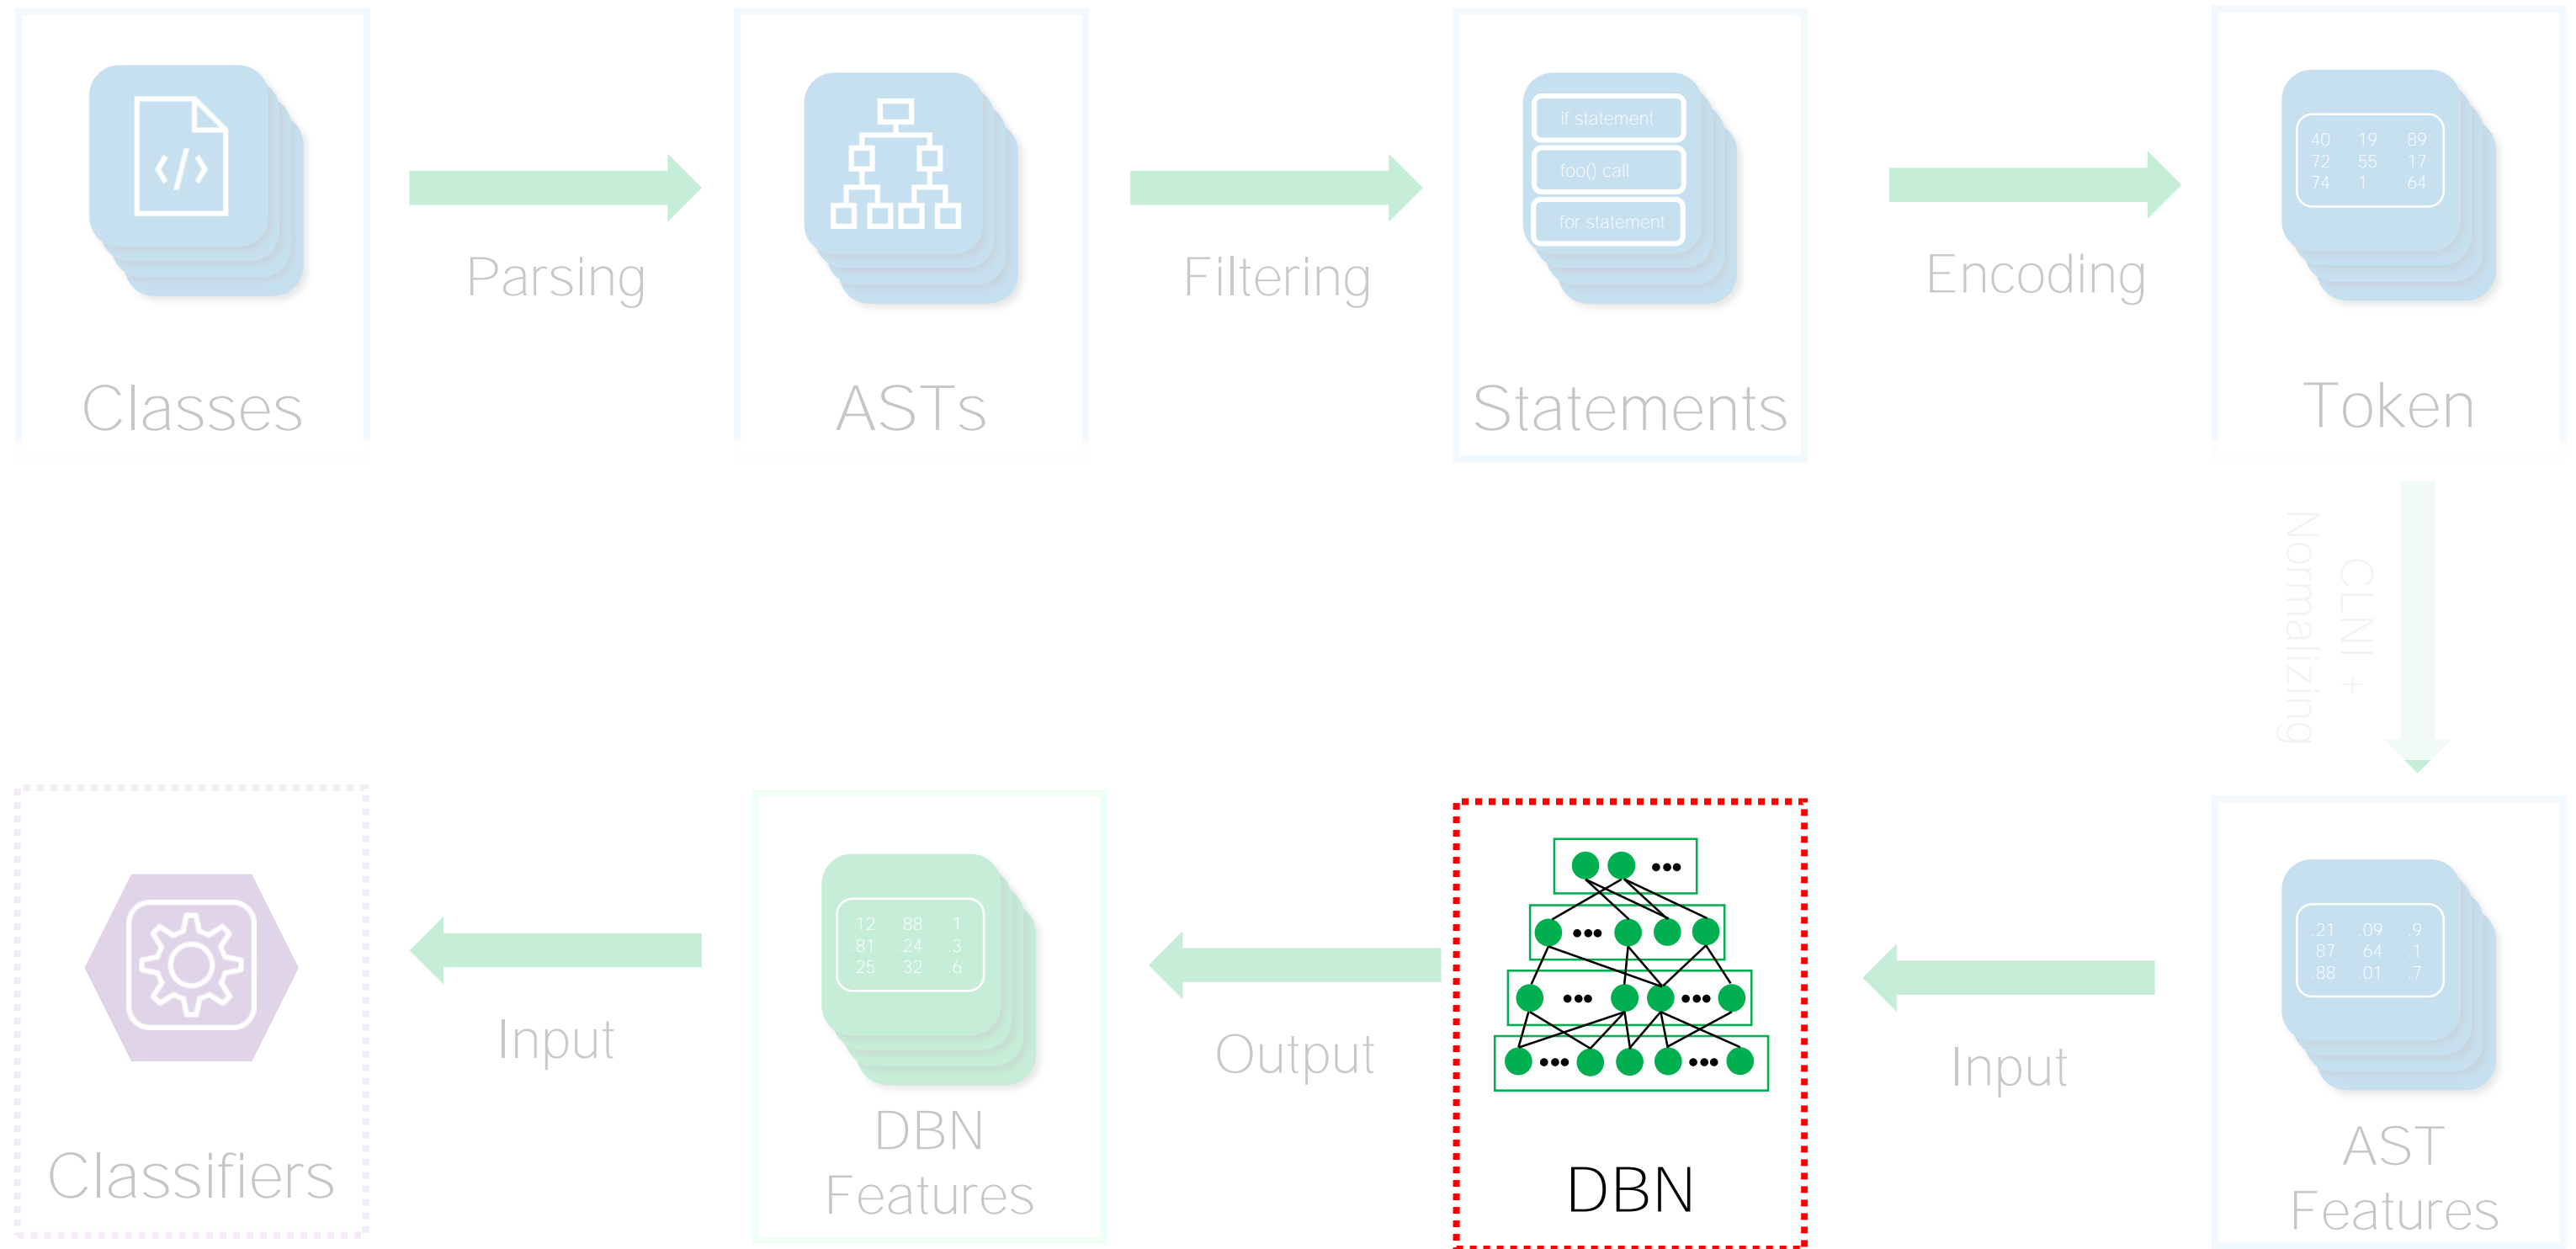


Conclusion

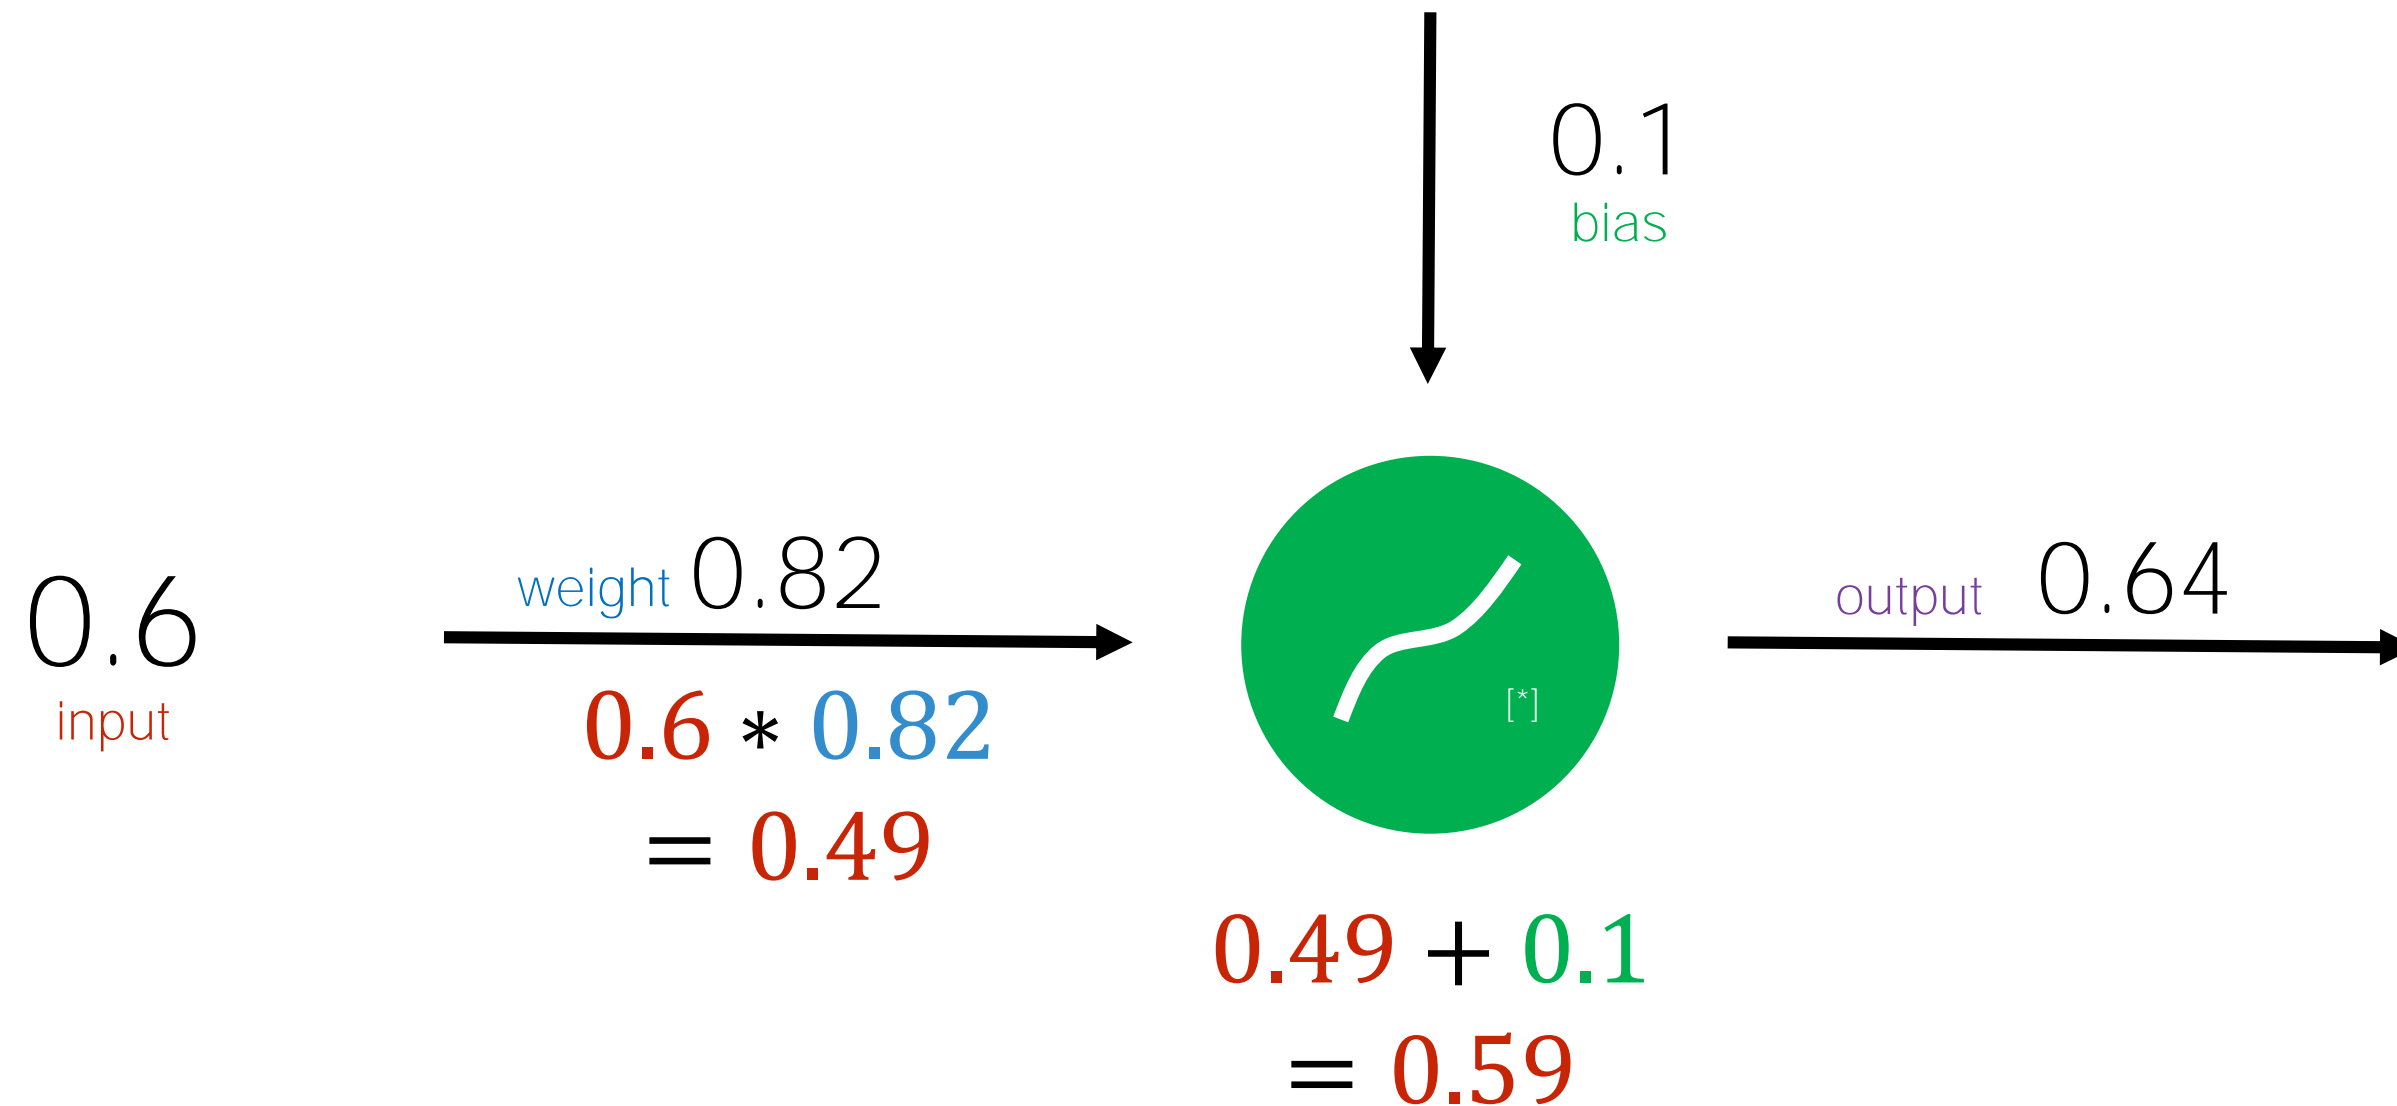
Prediction Process



Prediction Process




Recap: Perceptron




[*] Activation function. In this case:
Sigmoid function: $S(t) = \frac{1}{1+e^{-t}}$
In [8] this is **just a** „binary“ threshold.

Restricted Boltzmann Machine




Restricted Boltzmann Machine




- Shallow two layer network
- No intra layer connections (Restricted)
- Fully connected
- Can be trained unsupervised
- ‘Contrastive Divergence’ [9] / **‘Gibbs Sampling’** is used to train the network

Restricted Boltzmann Machine - Training




Sigmoid function: $S(t) = \frac{1}{1+e^{-t}}$

Restricted Boltzmann Machine - Training



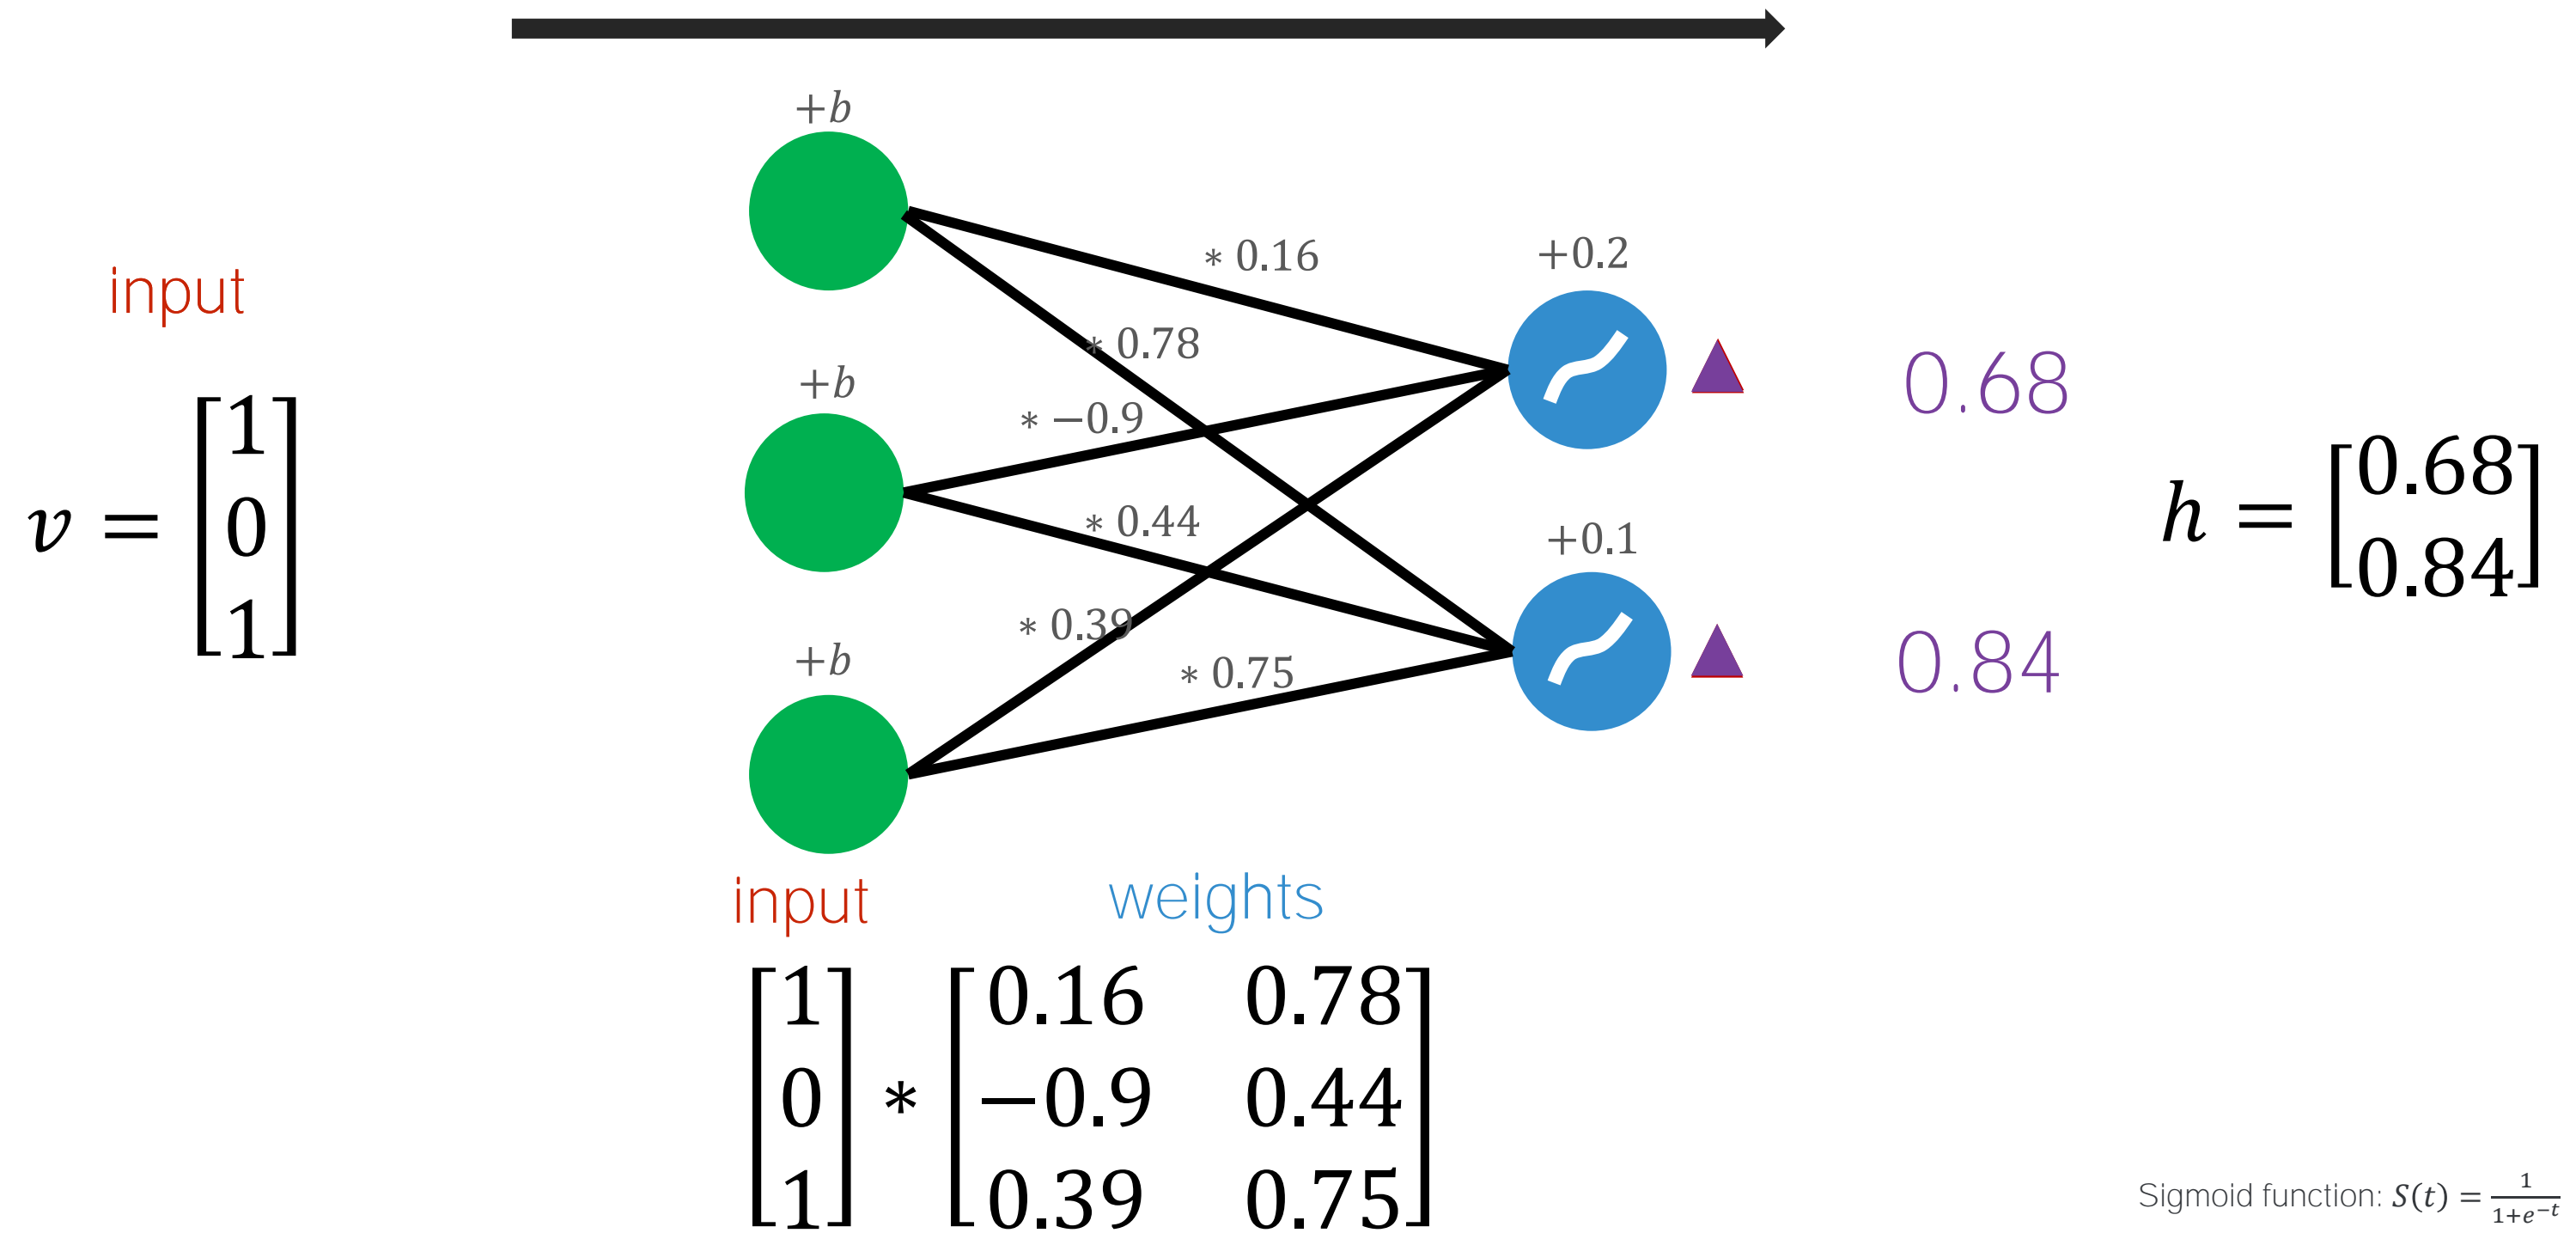
Sigmoid function: $S(t) = \frac{1}{1+e^{-t}}$

Restricted Boltzmann Machine - Training




Sigmoid function: $S(t) = \frac{1}{1+e^{-t}}$

Restricted Boltzmann Machine - Training



Restricted Boltzmann Machine - Training

$$v = \begin{bmatrix} 1 \\ 0 \\ 1 \end{bmatrix}$$




$$h = \begin{bmatrix} 0.68 \\ 0.84 \end{bmatrix}$$

$$\begin{bmatrix} 0.16 & -0.9 & 0.39 \\ 0.78 & 0.44 & 0.79 \end{bmatrix} * \begin{bmatrix} 1 \\ 1 \end{bmatrix}$$

Sigmoid function: $S(t) = \frac{1}{1+e^{-t}}$

Restricted Boltzmann Machine - Training

$$v = \begin{bmatrix} 1 \\ 0 \\ 1 \end{bmatrix}$$



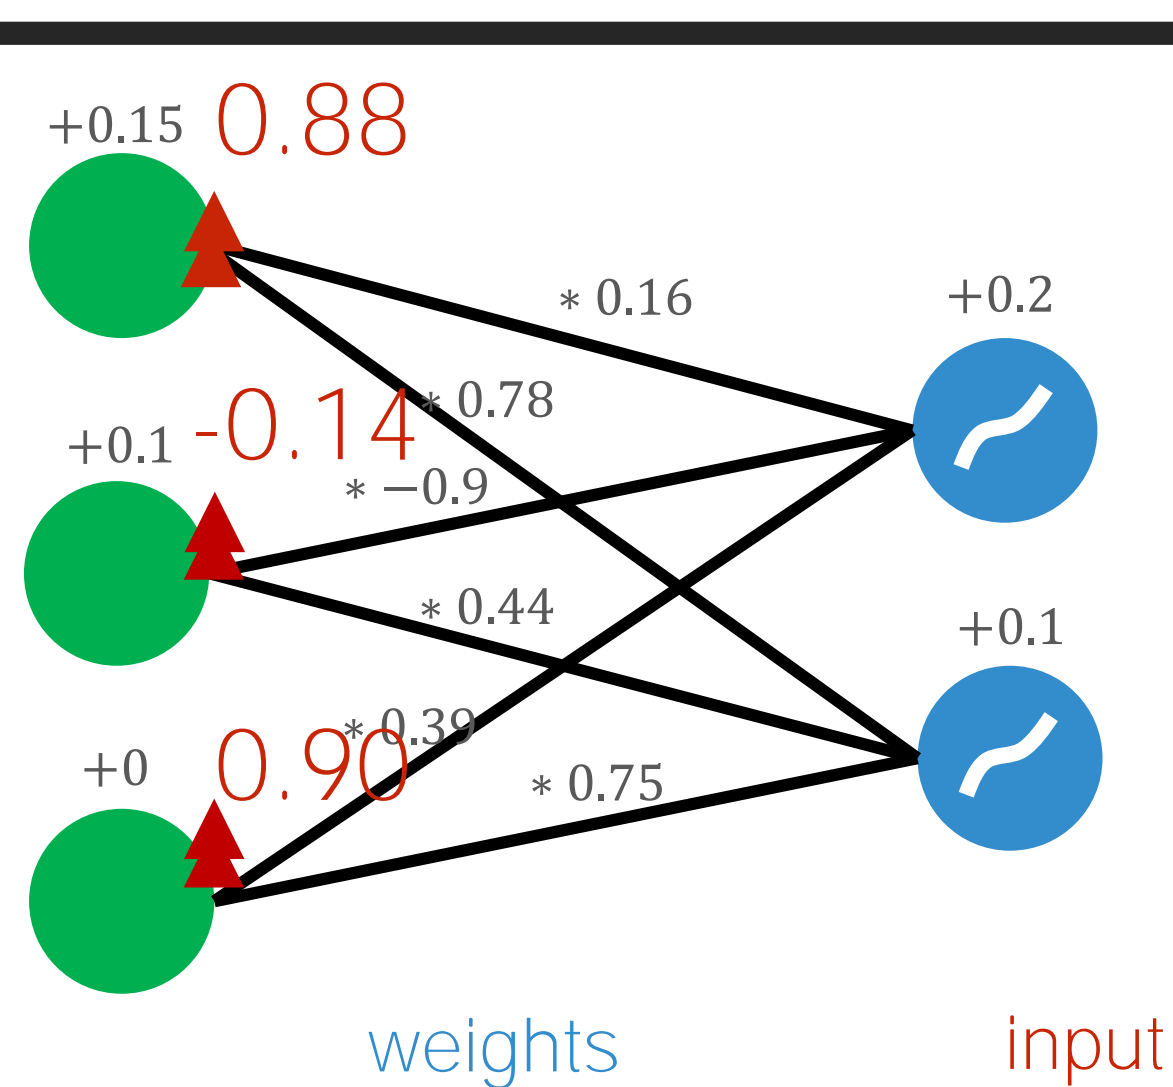
$$h = \begin{bmatrix} 0.68 \\ 0.84 \end{bmatrix}$$

$$\begin{bmatrix} 0.16 & -0.9 & 0.39 \\ 0.78 & 0.44 & 0.79 \end{bmatrix} * \begin{bmatrix} 1 \\ 1 \end{bmatrix}$$

$$\text{Sigmoid function: } S(t) = \frac{1}{1+e^{-t}}$$

Restricted Boltzmann Machine - Training

$$v = \begin{bmatrix} 1 \\ 0 \\ 1 \end{bmatrix}$$



$$h = \begin{bmatrix} 0.68 \\ 0.84 \end{bmatrix}$$


$$\begin{bmatrix} 0.16 & -0.9 & 0.39 \\ 0.78 & 0.44 & 0.79 \end{bmatrix} * \begin{bmatrix} 1 \\ 1 \end{bmatrix}$$

Sigmoid function: $S(t) = \frac{1}{1+e^{-t}}$

Restricted Boltzmann Machine - Training

$$v = \begin{bmatrix} 1 \\ 0 \\ 1 \end{bmatrix}$$

$$v' = \begin{bmatrix} 0.88 \\ -0.14 \\ 0.90 \end{bmatrix}$$



$$h = \begin{bmatrix} 0.68 \\ 0.84 \end{bmatrix}$$

$$\begin{bmatrix} 0.16 & -0.9 & 0.39 \\ 0.78 & 0.44 & 0.79 \end{bmatrix} * \begin{bmatrix} 1 \\ 1 \end{bmatrix}$$


Sigmoid function: $S(t) = \frac{1}{1+e^{-t}}$

Restricted Boltzmann Machine - Training

$$v = \begin{bmatrix} 1 \\ 0 \\ 1 \end{bmatrix}$$

$$v' = \begin{bmatrix} 0.88 \\ -0.14 \\ 0.90 \end{bmatrix}$$

Calculate error and
adjust weights so
that $v \approx v'$



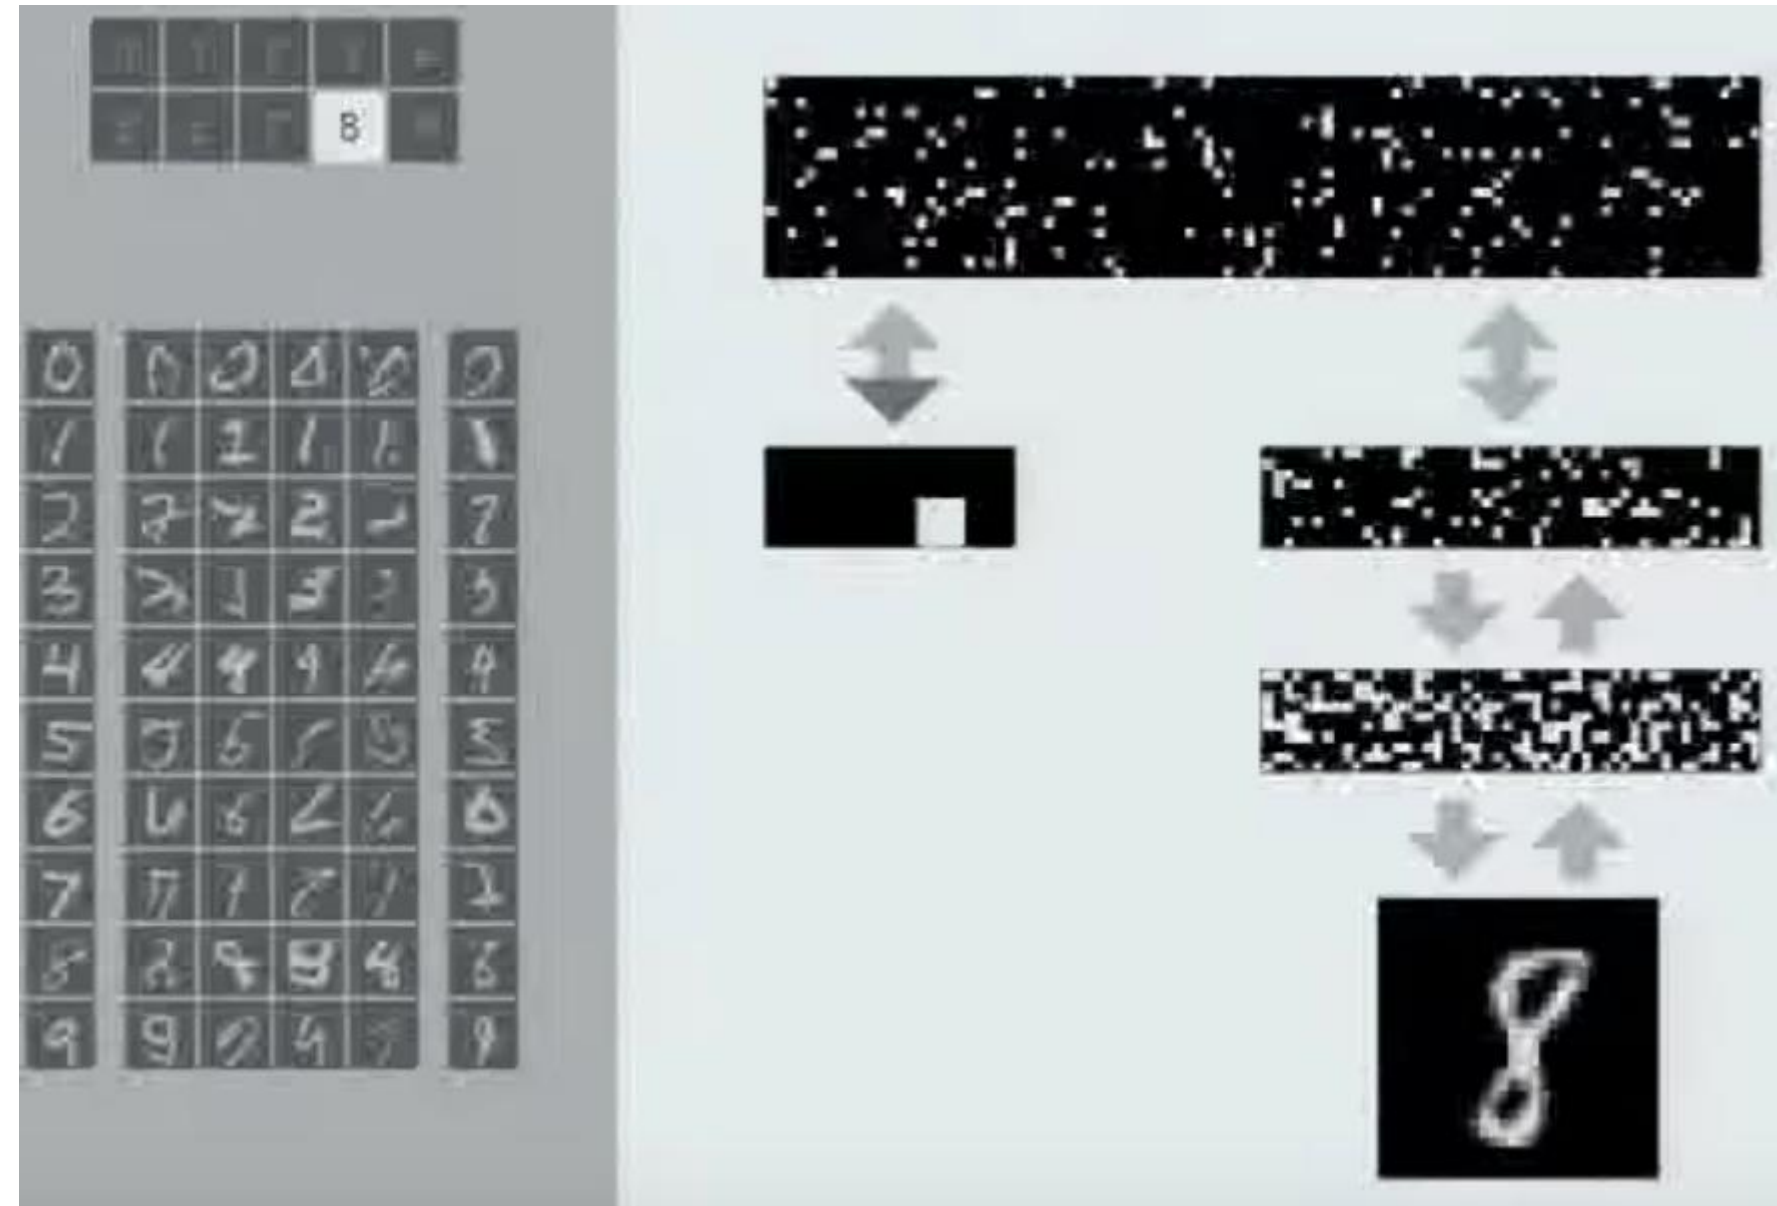
Sigmoid function: $S(t) = \frac{1}{1+e^{-t}}$

Restricted Boltzmann Machine - Training


$$v = \begin{bmatrix} 1 \\ 0 \\ 1 \end{bmatrix}$$

$$v' = \begin{bmatrix} 0.88 \\ -0.14 \\ 0.90 \end{bmatrix}$$

Calculate error and
adjust weights so
that $v \approx v'$




Restricted Boltzmann Machine - Training




How often are they on together with input v ?

Restricted Boltzmann Machine - Training




How often are they on together when the model is dreaming?

Restricted Boltzmann Machine - Training




Restricted Boltzmann Machine - Training




This will force towards „dreaming“ the actual data


Deep Belief Network




Deep Belief Network




Deep Belief Network



Deep Belief Network




Prediction Process



Agenda




Defect Prediction



Technical Background



Discussion of Results



Demo



Conclusion

Results – Within Project Defect Prediction


- Success is measured with F1-Score
- Baselines:



Code
metrics

20 traditional features:

- Lines of code
- Operand count
- McCabe complexity



AST
Features


Filtered AST features

Results – Within Project Defect Prediction

- Success is measured with F1-Score
- Baselines:
 - 20 traditional features
 - AST features
- Classifiers:
 - ADTree
 - Naïve Bayes
 - Logistic Regression

Results – Within Project Defect Prediction


Results for Within Project Defect Prediction using ADTrees for Classification



[11] S. Wang, T. Liu, L. Tan. "Automatically learning semantic features for defect prediction." In: Proceedings - International Conference on Software Engineering (2016)


Results – Within Project Defect Prediction

Results for Within Project Defect Prediction using ADTrees for Classification



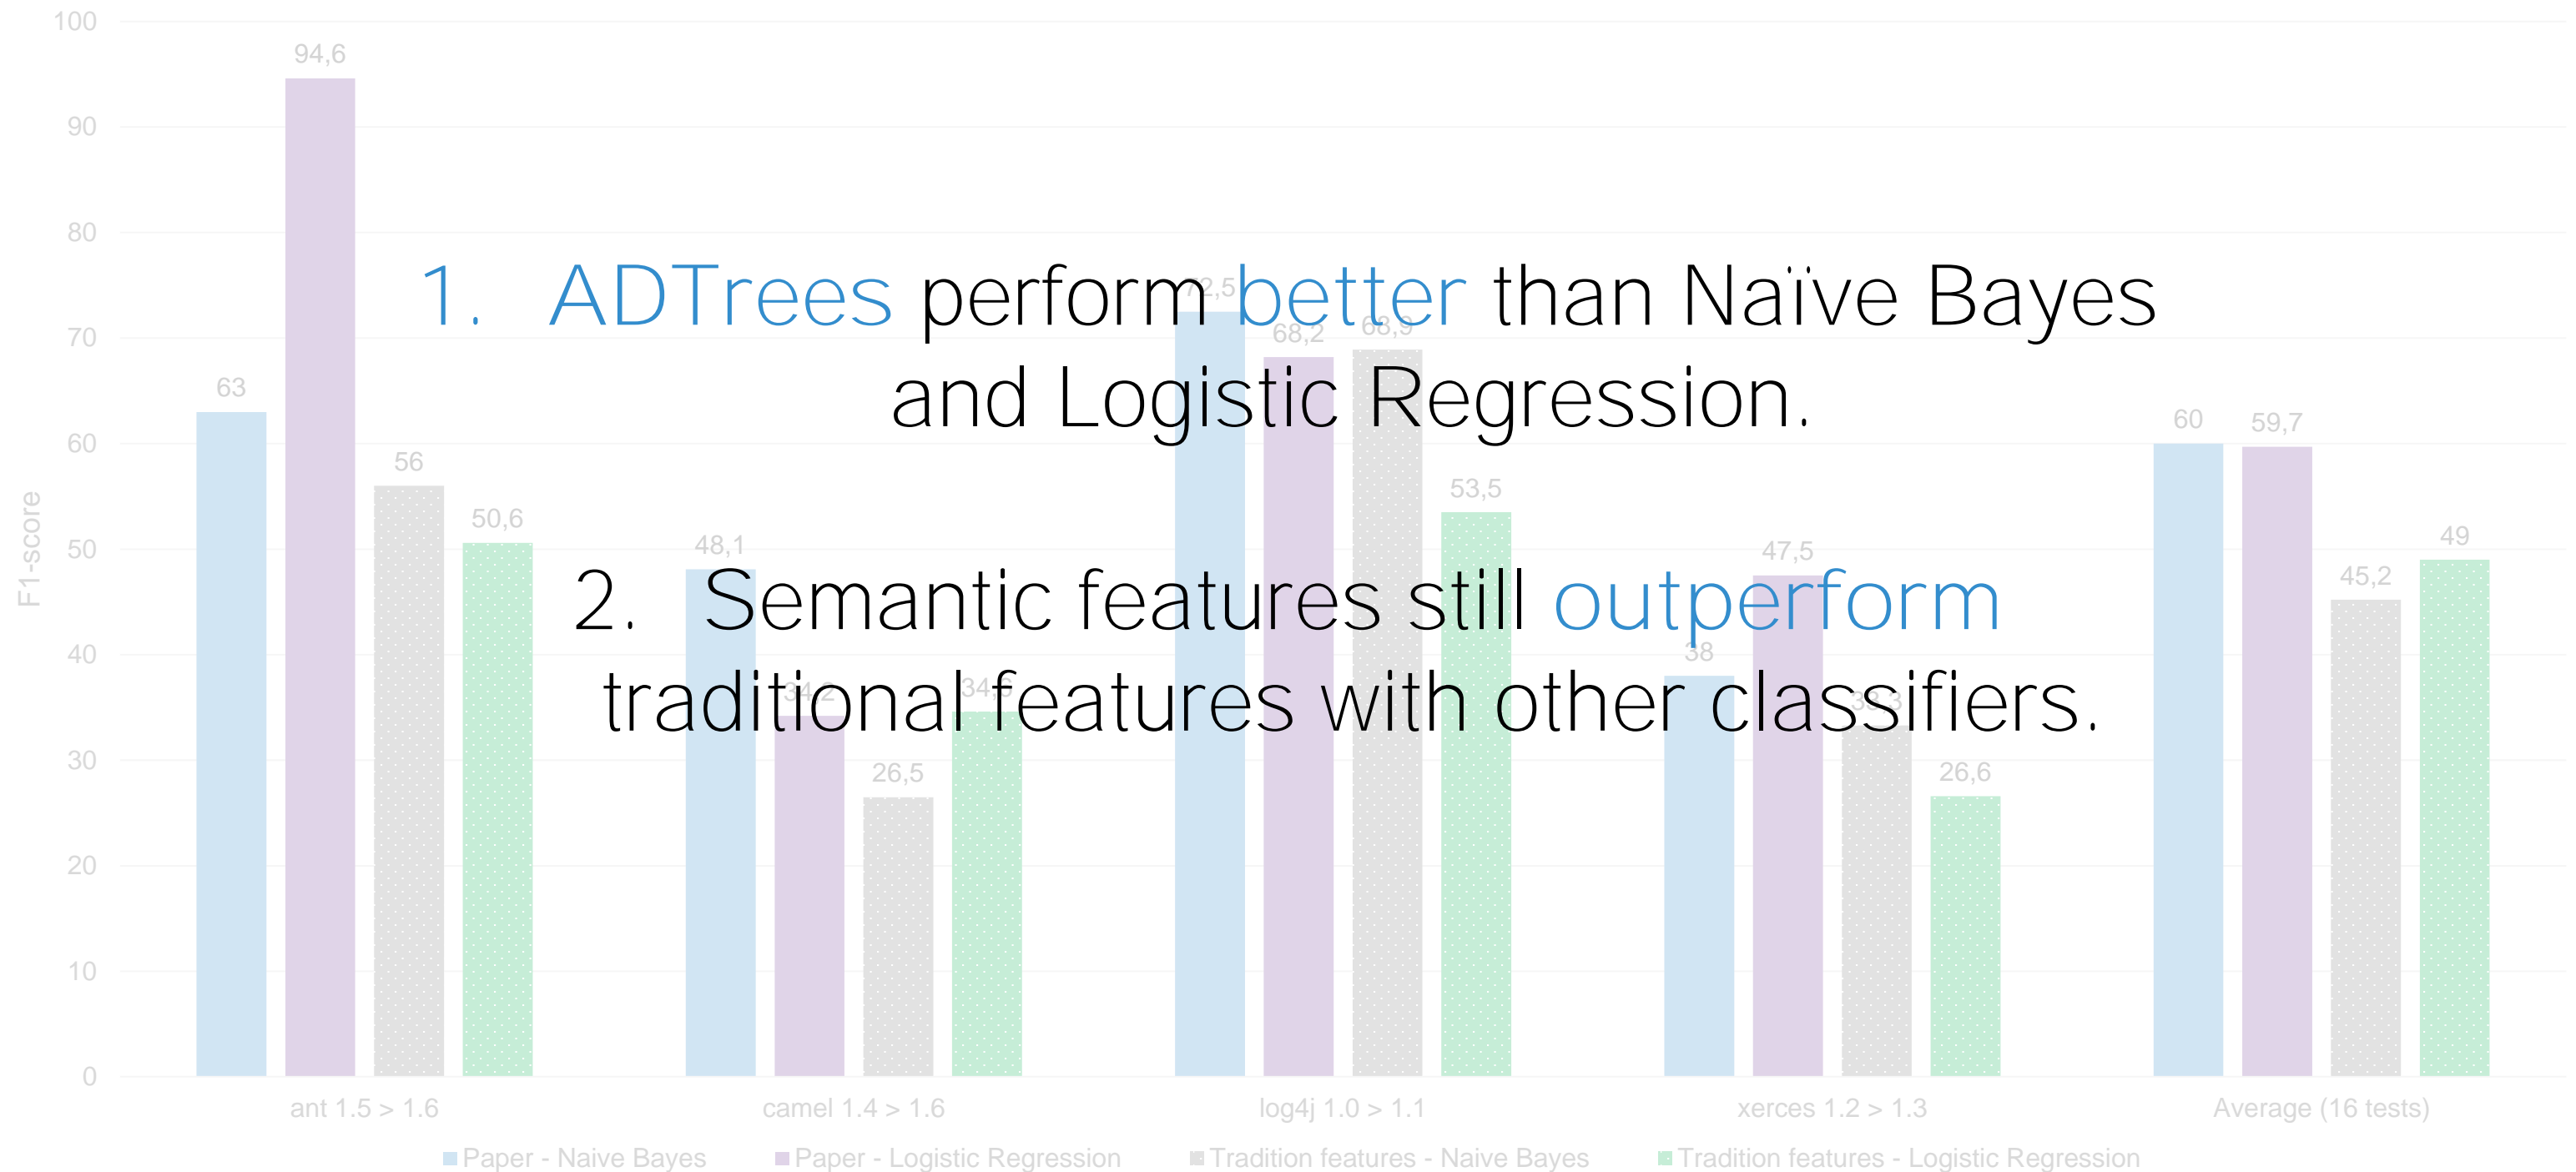
Results – Within Project Defect Prediction

Results for Within Project Defect Prediction using Different Classifiers




Results – Within Project Defect Prediction

Results for Within Project Defect Prediction using Different Classifiers




Results – Cross Project Defect Prediction

Results for Cross Project Defect Prediction using ADTrees and TCA+ for Classification




Results – Cross Project Defect Prediction

Results for Within Project Defect Prediction using ADTrees for Classification




Results – Cross Project Defect Prediction

Results for Within Project Defect Prediction using ADTrees for Classification




Results – Cross Project Defect Prediction

Results for Within Project Defect Prediction using ADTrees for Classification



Results – Cross Project Defect Prediction


Results for Within Project Defect Prediction using ADTrees for Classification



Results – How could the results be improved?

Results for Within Project Defect Prediction using ADTrees for Classification

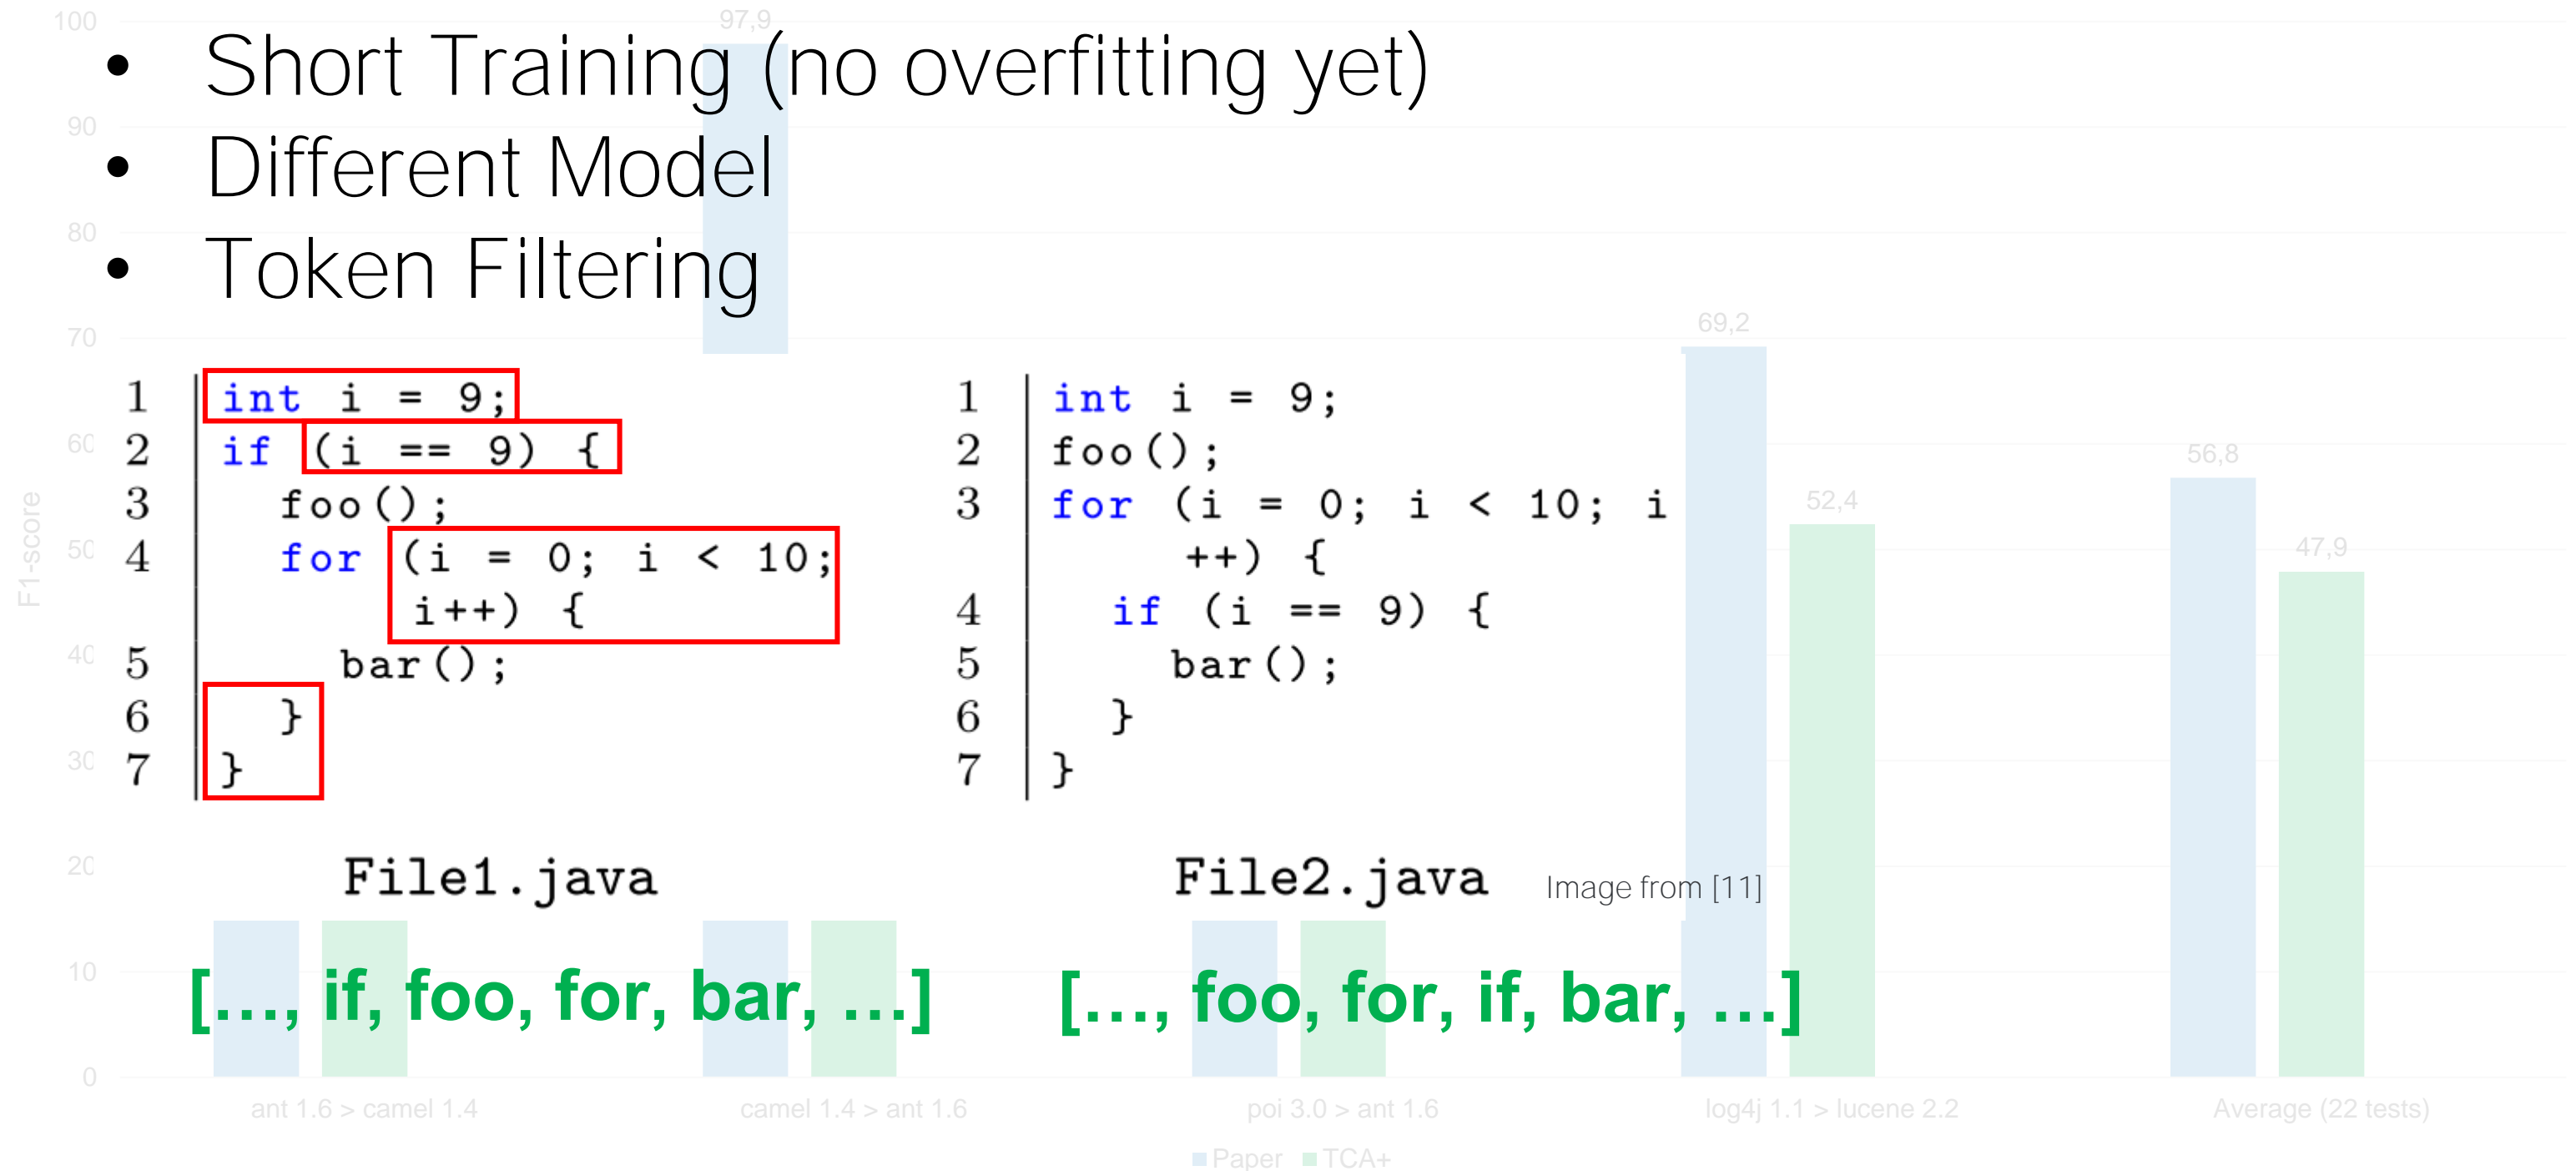
- Short Training (no overfitting yet)



Results – How could the results be improved?

Results for Within Project Defect Prediction using ADTrees for Classification


- Short Training (no overfitting yet)
- Different Model
- Token Filtering



Agenda



Defect Prediction



Technical Background



Discussion of Results



Demo



Conclusion

Defect Prediction - Demo



Can **normal Neural Networks** with GPU-Training improve the F1-Score?




Does the F1-Score increase if information about the **control flow** of the program is included?

Defect Prediction - Demo



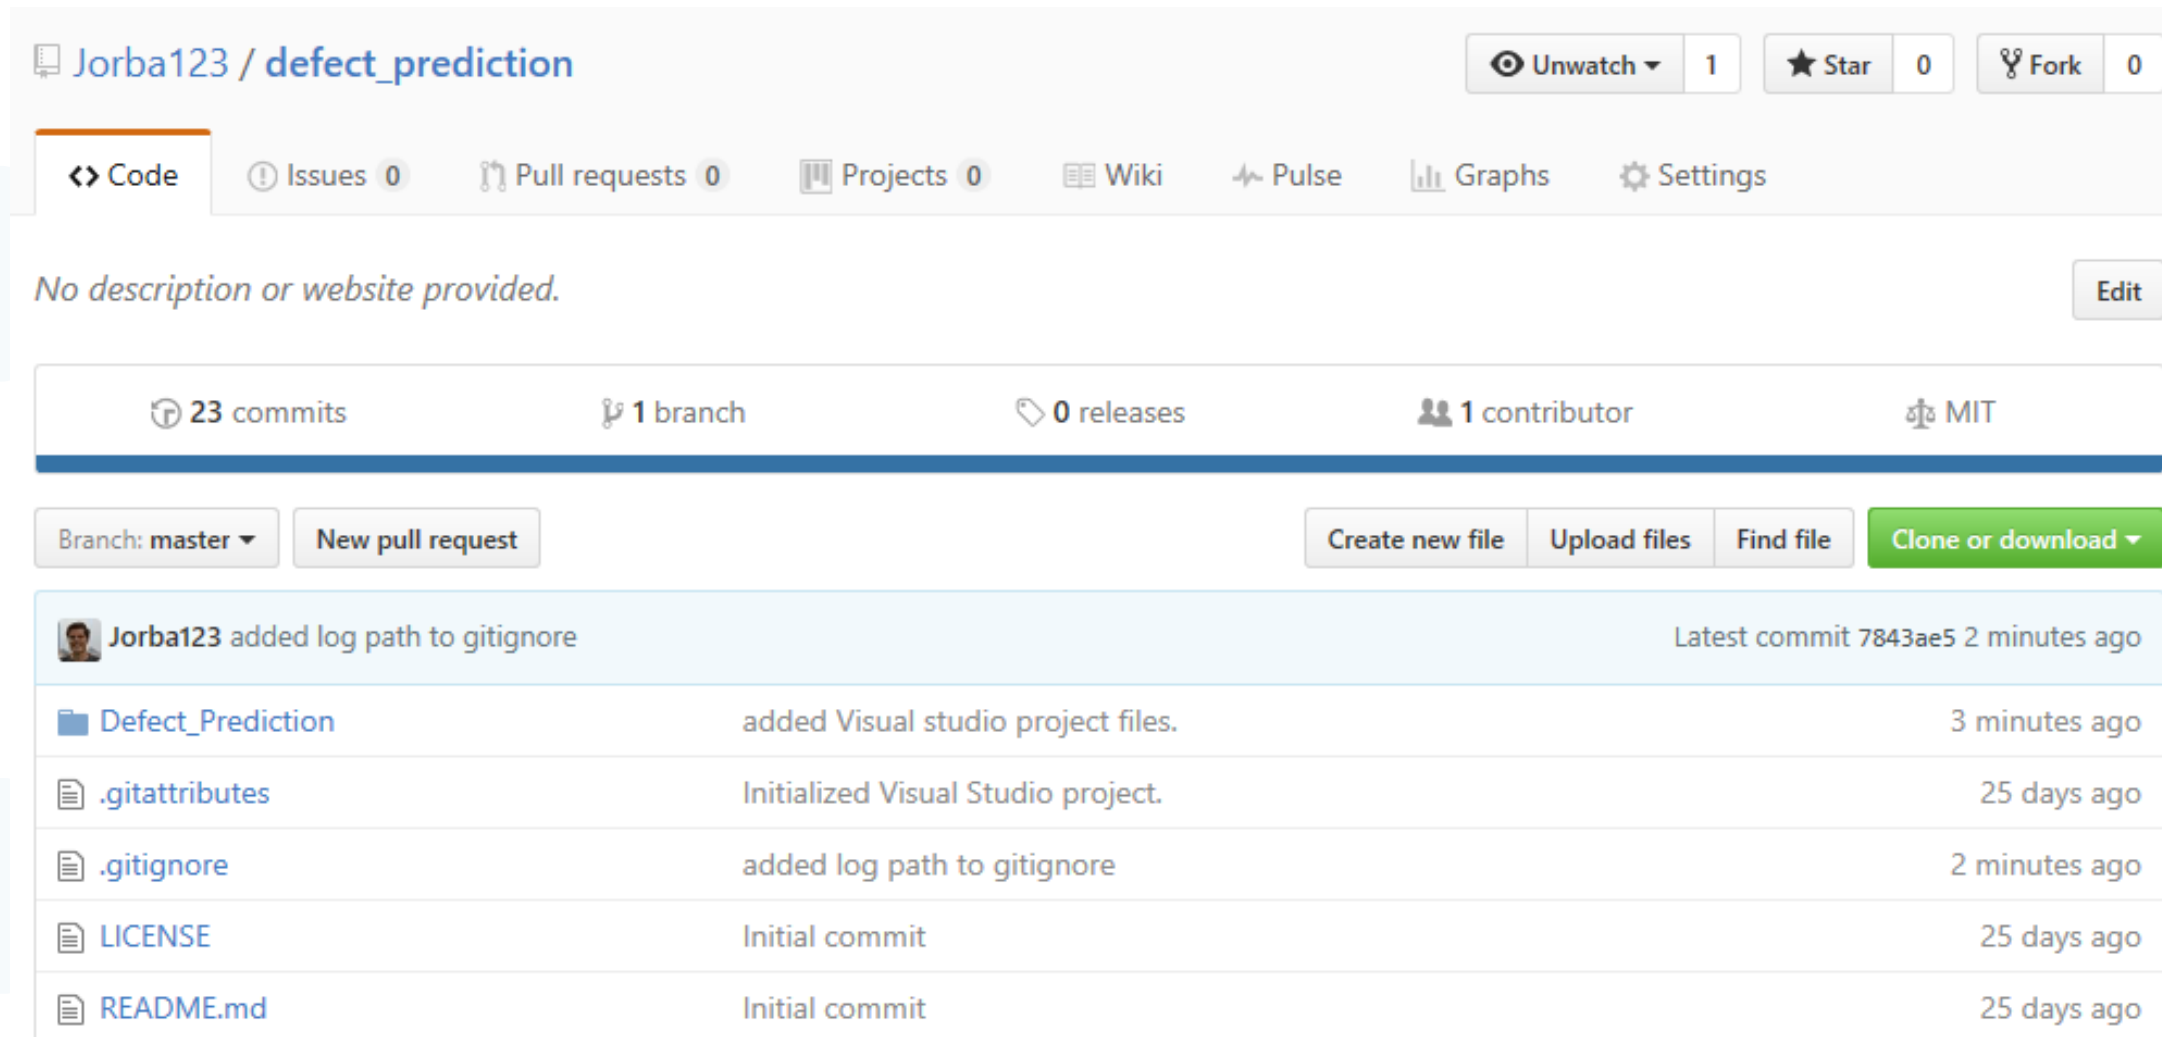
Can normal Neural Networks with GPU-
Training improve the F1-Score?

NO 

Does the F1-Score increase if information
about the control flow of the program is
included?

(but I will show you the program nonetheless)

Defect Prediction - Demo



The screenshot shows a GitHub repository page for 'Jorba123 / defect_prediction'. The repository has 1 watch, 0 stars, and 0 forks. The 'Code' tab is selected, showing a message 'No description or website provided.' and an 'Edit' button. Below this, statistics show 23 commits, 1 branch, 0 releases, 1 contributor, and the MIT license. Action buttons include 'Branch: master', 'New pull request', 'Create new file', 'Upload files', 'Find file', and 'Clone or download'. A commit history table is displayed below.

Commit Message	Time Ago
Jorba123 added log path to gitignore	Latest commit 7843ae5 2 minutes ago
Defect_Prediction added Visual studio project files.	3 minutes ago
.gitattributes Initialized Visual Studio project.	25 days ago
.gitignore added log path to gitignore	2 minutes ago
LICENSE Initial commit	25 days ago
README.md Initial commit	25 days ago

Download or clone at
https://github.com/Jorba123/defect_prediction

Discussion of Demo Result



Number of **training samples** is very small



Data is **imbalanced** (subsampling not possible)



Sparse input features due to padding



Possible **labeling errors** in data set




Class-wise bug labeling

Agenda




Defect Prediction



Technical Background



Discussion of Results




Demo



Conclusion

Conclusion



Defect Prediction can **lower cost** and **save time** during and after development



Deep Belief Networks (DBNs) significantly **outperform traditional features** and classifiers



Defect prediction rates differ a lot across projects



Prediction rates for cross project prediction are lower than for within project prediction

Thank you!



References

- [1] F. Brady “Cambridge University Study States Software Bugs Cost Economy \$312 Billion Per Year” prweb (2013) accessed 12.01.2017 <http://www.prweb.com/releases/2013/1/prweb10298185.htm>
- [2] R. Pielke, R. Byerly “Shuttle Programme Lifetime Cost.” In Nature Vol. 472 (2011)
- [3] “Milliardengrab A380” Handelsblatt (2012) <http://www.genios.de/presse-archiv/artikel/HB/20121108/milliardengrab-a380/2B70894B-6C35-4D7B-A956-B0B5783C418D.html>
- [4] World Bank – Last accessed 12.01.2017 http://data.worldbank.org/indicator/NY.GDP.MKTP.CD?locations=GR&name_desc=true
- [5] S. McConnell “Code Complete 2nd edition.” Microsoft Press (2004) page 521
- [6] S. Lohr, J. Markoff “Windows Is So Slow, but Why?” New York Times (2006) Last accessed 12.01.2017 <http://www.nytimes.com/2006/03/27/technology/27soft.html?adxnnl=1&pagewanted=all&adxnnlx=1382805118-0jnNRGXEPip3xoW+BDp8Q>
- [7] B. Boehm, P. Papaccio “Understanding and Controlling Software Costs.” In: IEEE Transactions on Software Engineering (1988) page 1466
- [8] D. Ackley, G. Hinton, T. Sejnowski “A learning algorithm for boltzmann machines.” In: Cognitive Science (1985) Volume 9

References

- [9] G. Hinton, S. Osindero, Y. Teh “A Fast Learning Algorithm for Deep Belief Nets.” In: Neural Computation (2006) Volume 18
- [10] G. Hinton. “The Next Generation of Neural Networks.” Google TechTalks (2007)
<https://www.youtube.com/watch?v=AyzOUbkUf3M> Last accessed: 12.01.2017
- [11] S. Wang, T. Liu, L. Tan. “Automatically learning semantic features for defect prediction.” In: Proceedings - International Conference on Software Engineering (2016)
- All Icons from Icons8 <https://icons8.com/>